

RYAN TRESOR KEMAJOU

ryantresor1.k@hotmail.com | 240-733-2116 | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

Data scientist with nearly four years of experience turning operational data into decisions, currently completing an M.S. in Data Science at the University of Maryland. Builds and measures end-to-end machine learning systems, covering annotation, baselines, trained models, cross-validated evaluation, and deployment, with a record of reporting results honestly, including where they fall short.

EDUCATION

University of Maryland

Master of Science in Data Science

Bachelor of Science in Computer Science, Data Science

College Park, MD

Sept 2025 – Dec 2026

Aug 2023 – May 2025

Coursework: Natural Language Processing, Machine Learning, Causal Inference, Bayesian Networks, Optimization & Linear Programming, Big Data Systems, Probability & Statistics, Linear Algebra.

Leadership: ColorStack UMD (Secretary), Code Black member. **Certification:** CompTIA Security+.

TECHNICAL SKILLS

Languages: Python, SQL, R, TypeScript/JavaScript, Java, C++

ML & NLP: scikit-learn, PyTorch, XGBoost, logistic regression, random forests, cross-validation, annotation & inter-annotator agreement, LLM evaluation, prompt engineering

Data & Infra: PostgreSQL, pandas, NumPy, Flask, FastAPI, Streamlit, Docker, AWS, Celery/Redis, Git, Tableau

EXPERIENCE

Ylunga Consulting

Founder & Data/ML Consultant

Montgomery County, MD

Jan 2026 – Present

- Ship production data and ML tools for healthcare and public-sector clients, owning the full path from problem framing through deployment, billing, and support.
- Built a clinical documentation platform on Flask and PostgreSQL serving LLM inference through a server-side proxy, with CSRF protection, rate limiting, security headers, and a 15-minute HIPAA session timeout.
- Built a bilingual procurement scout that scrapes six government sources across two countries and ranks each opportunity 0–10 with an LLM scorer, using separate French and English prompts and prompt caching to cut per-run cost.

Abundant Grace Health Services

Data Analyst

Washington, DC

Oct 2021 – Aug 2025

- Analyzed operational and client datasets in Python and SQL to surface the utilization and staffing trends that informed executive planning.
- Built and maintained automated KPI dashboards in Tableau, replacing a manual weekly reporting cycle and giving leadership a standing view of performance.
- Partnered with healthcare providers to define risk-based metrics and implemented a classification-based scoring tool associated with an 18% improvement in client outcomes.
- Automated recurring reporting and data-cleaning pipelines, reducing manual analysis time by 60%.
- Presented findings to senior management, translating model output into decisions non-technical stakeholders could act on.

PROJECTS

NLAH-Builder: NLP research system | [GitHub](#)

Python, scikit-learn, Streamlit

- Led data and evaluation for a 3-person NLP team, contributing 82 of 113 commits: hand-labeled 30 real-world documents into **763 gold slot values**, wrote the annotation guidelines, and ran an intra-annotator agreement study.
- Replaced a keyword-rule extractor with candidate generation plus a logistic-regression classifier, raising micro F1 from **0.05 to 0.20** and precision from 0.17 to 0.36 under leave-one-document-out cross-validation.
- Established the ceiling rather than claiming the win: only 42% of gold values appear verbatim in source, and a blind second labeling pass scores 0.30, placing the model at **55% of human performance**.
- Designed a 20-trial blind downstream evaluation across four conditions; the generated rules file cut forbidden-file reads from 3 of 3 trials to **0 of 3**.

Denial Report Builder: healthcare revenue-cycle analytics

Python, Streamlit, SQLAlchemy, Stripe, Docker

- Built a 2,100-line application that ingests payer denial spreadsheets and returns a ranked, category-mapped report, with authentication, subscription gating, and a Stripe webhook service.